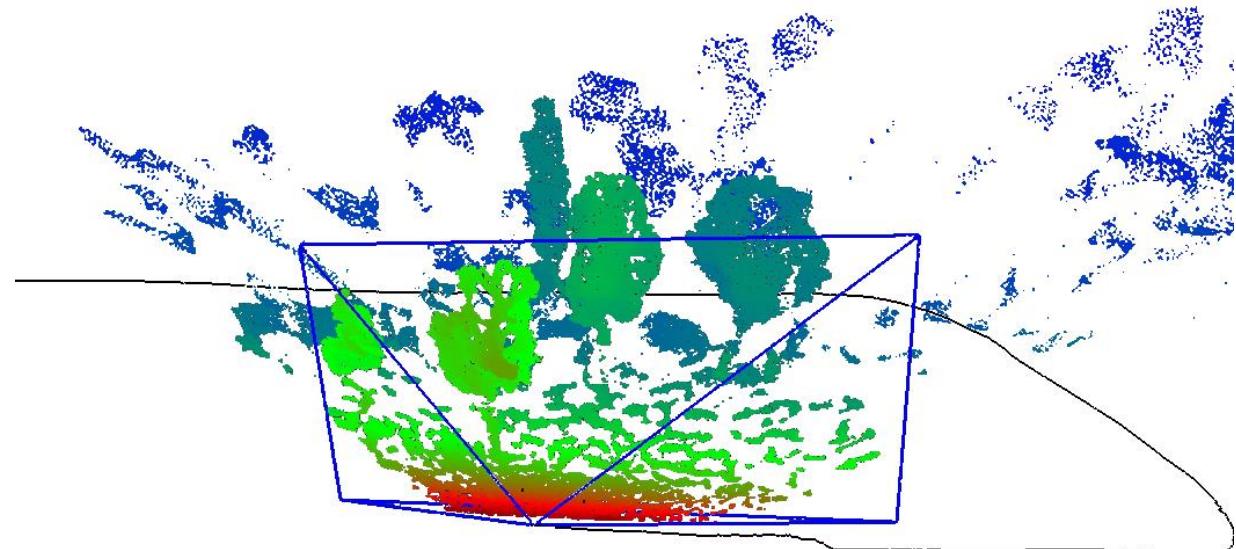




Semi-Dense Visual Odometry for a Monocular Camera

Jakob Engel, Jürgen Sturm, Daniel Cremers

Intl. Conf. on Computer Vision (ICCV) 2013



Monocular Video

Camera Motion and Scene Geometry

Visual Odometry



Camera: 752x480 @ 30fps, global shutter, monochrome, 130° diagonal fov

Applications

Augmented / Virtual Reality



- requires **camera pose** to render objects
- requires **scene geometry** e.g. for physical interaction

Robotics

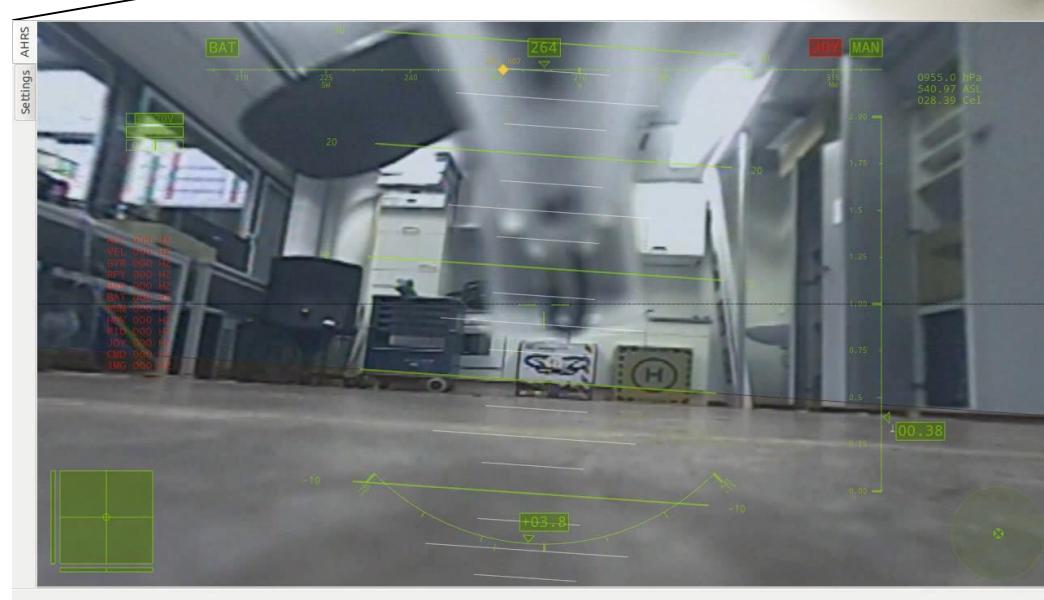


- requires **camera pose** to control robot position
- requires **scene geometry** e.g. to avoid obstacles

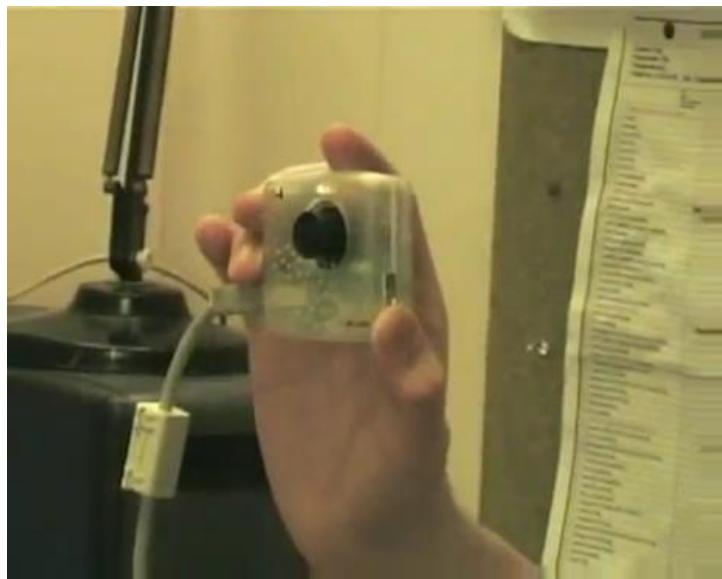


Why Monocular?

- small, light-weight
 - cheap
 - low power comsumption
 - versatile (scale-ambivalent)
 - easy(er) to calibrate



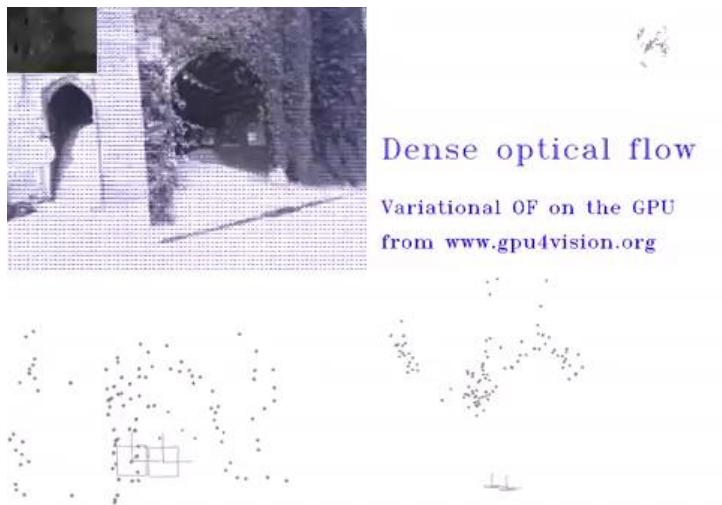
Long History...



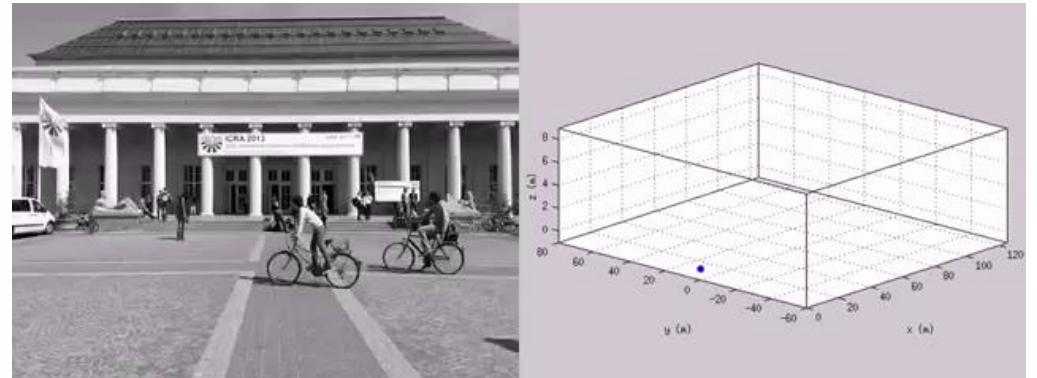
A. Davison et.al., 2004



G. Klein et.al., 2007



H. Strasdat et.al., 2010



A. Mourikis et.al., 2013

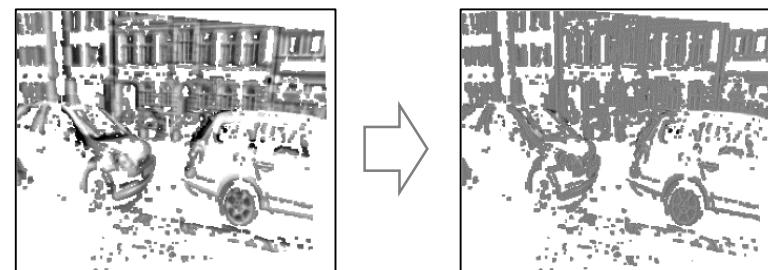
Overview

Input Video 640x480 @ 30Hz



Tracking

Estimate pose of new frame relative to depth map

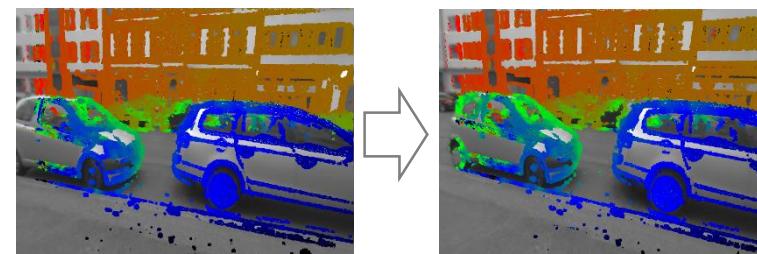


Separate Initialization

- 8 point algorithm
- random

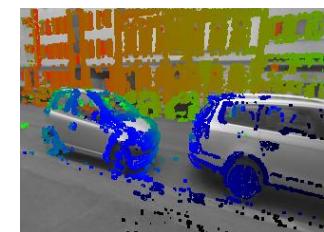
Mapping

Propagate and update depth map (stereo)



Semi-Dense Depth Map

- pixel depth
- depth variance

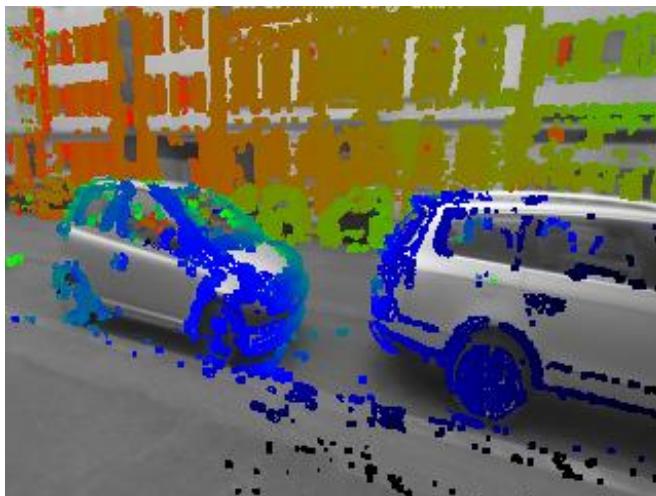




- image



World Representation



inv. depth mean



inv. depth variance



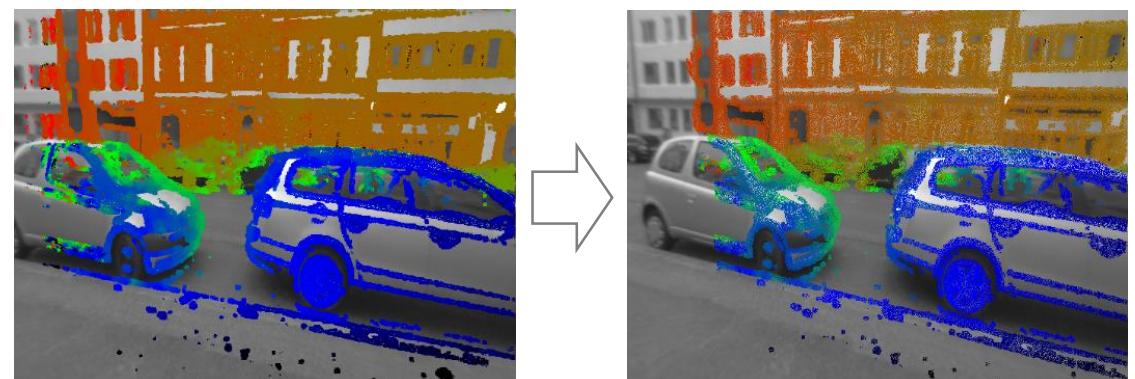
intensity image

- Gaussian on inverse depth
- for all pixel with sufficient abs. image gradient
- pose-graph + keyframes



Mapping

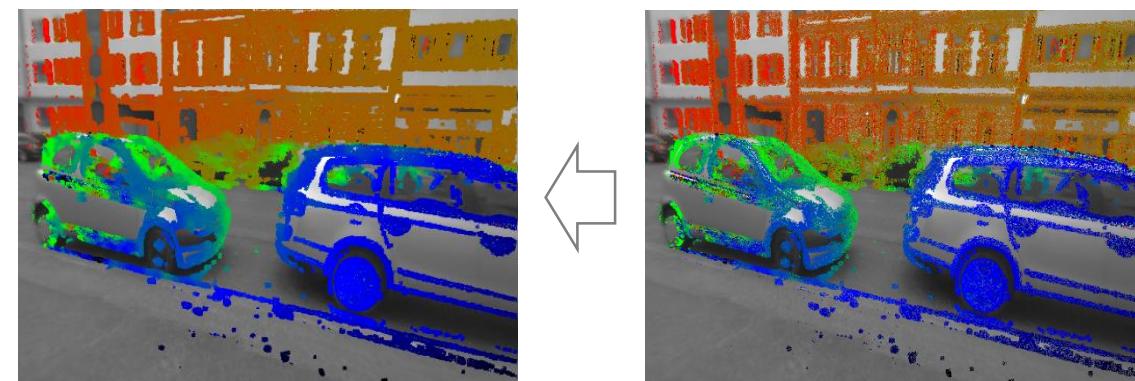
1. Propagation



2. Stereo-update

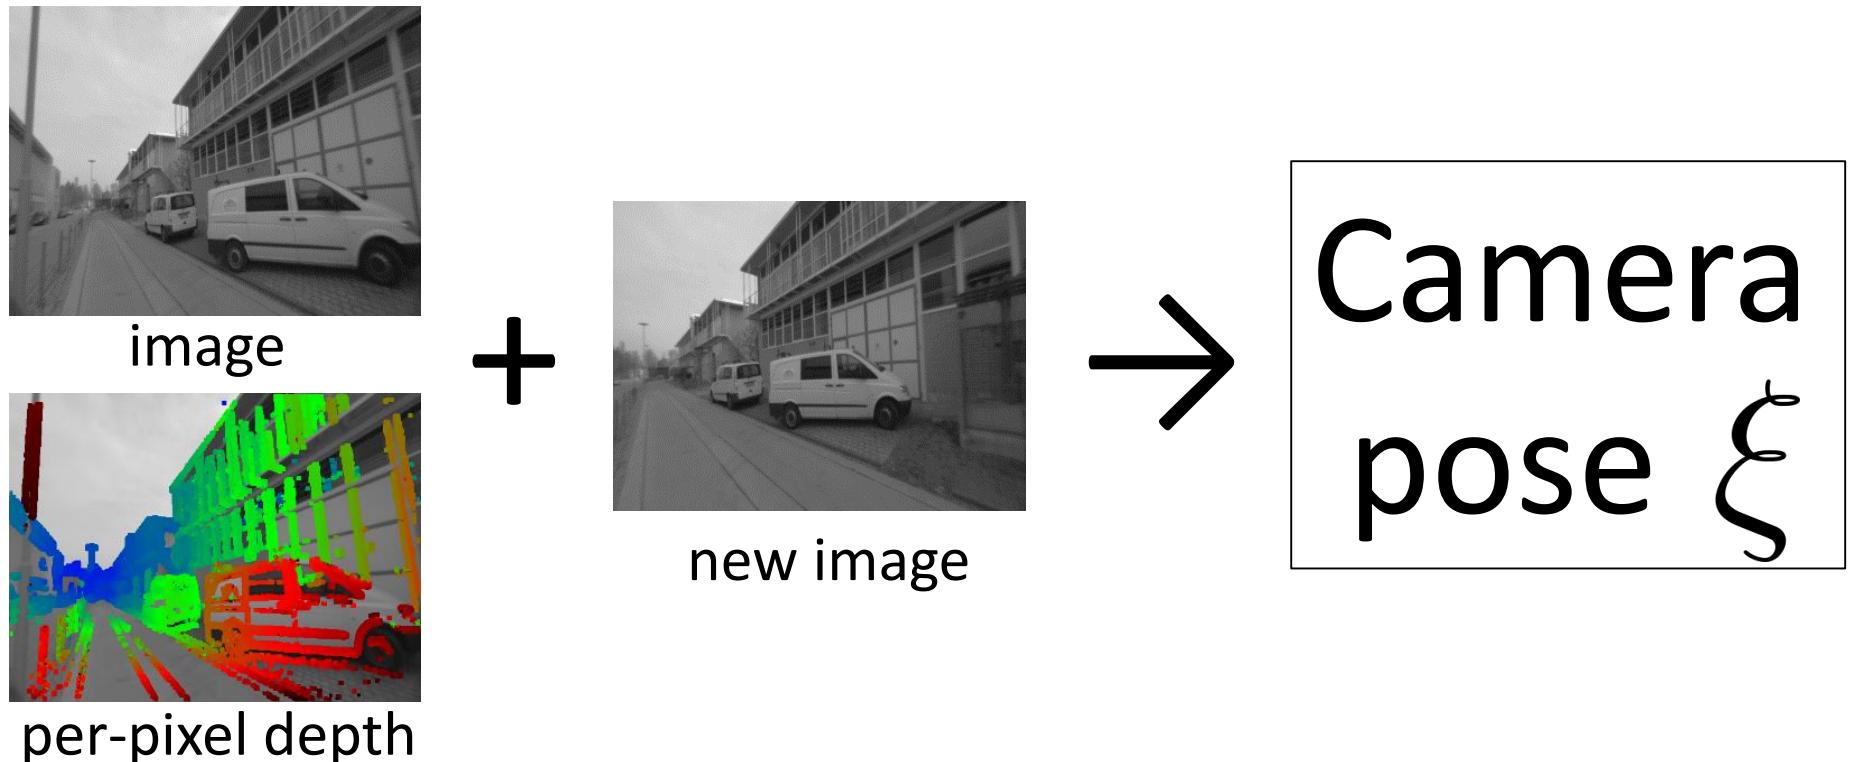


3. Regularization



Direct Image Alignment

- = „Direct Tracking“ / „Dense Tracking“
- = „Lucas-Kanade Tracking on $SE(3)$ “
- Maximum-Likelihood Estimator
- often used for RGB-D tracking (Kinect)
(Kerl et.al. @ ICRA '13; Steinbruecker et.al. @ ICCV '11; and many more)



Direct Image Alignment

Direct minimization of photometric error

$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} (I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2$$

camera pose ref. image new image ref. depth
 sum over valid ref. pixel

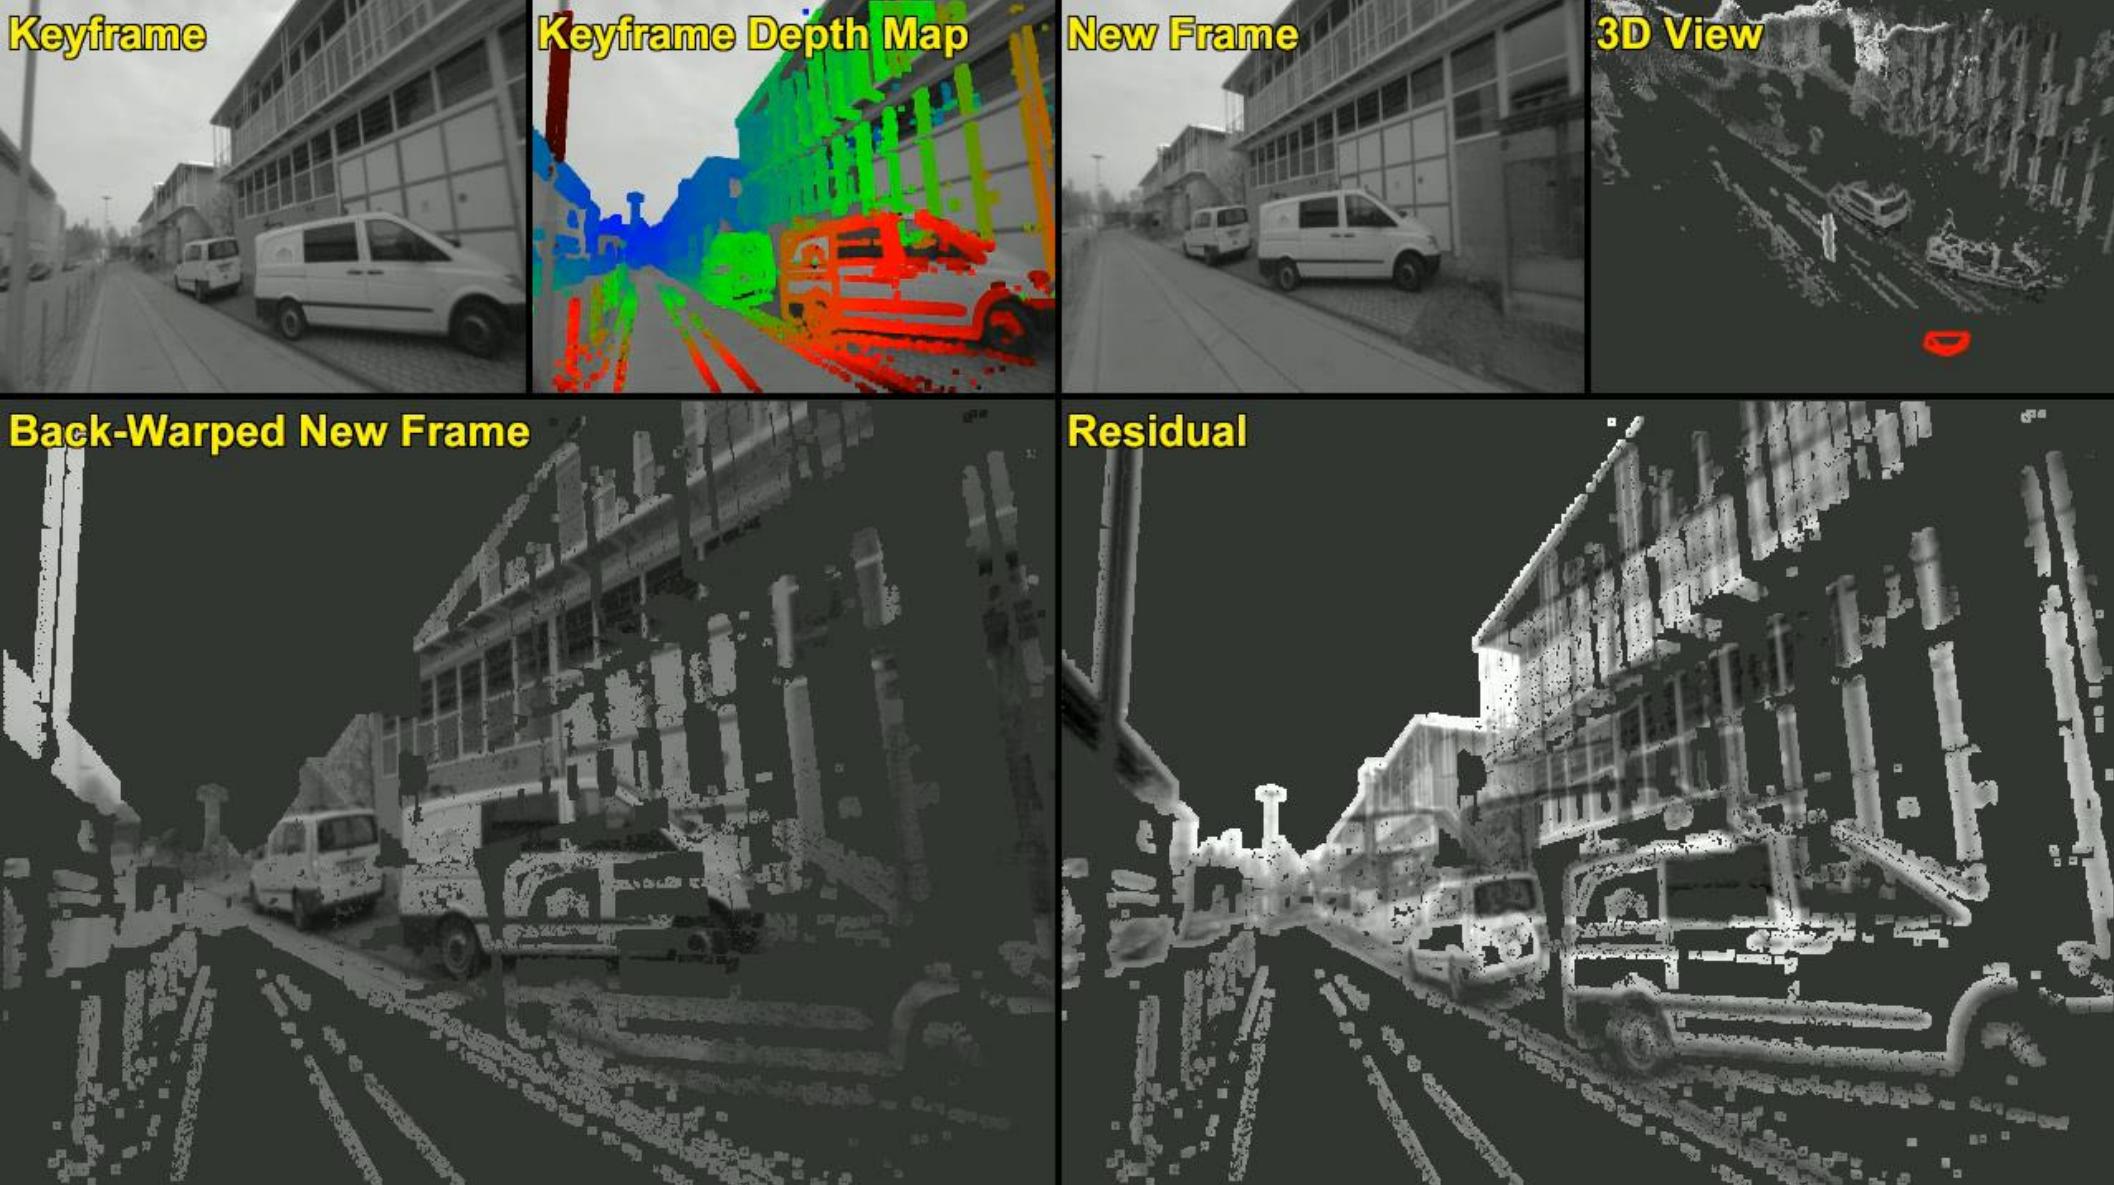
$$\omega(\mathbf{p}_i, d, \xi) := \pi(K(R_\xi K^{-1} \begin{pmatrix} dp_{i,x} \\ dp_{i,y} \\ d \end{pmatrix} + \mathbf{t}_\xi))$$

$$\pi(x, y, z) := \begin{pmatrix} x/z \\ y/z \end{pmatrix}$$

$$\begin{pmatrix} R_\xi & \mathbf{t}_\xi \\ 0 & 1 \end{pmatrix} := \exp(\hat{\xi})$$

$\omega(\mathbf{p}_i, d, \xi)$ „warps“ a pixel from ref. image to new image

Direct Image Alignment



$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} (I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2$$

$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} (I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2$$

- solved using the **Gauss-Newton** algorithm using left-multiplicative increments on $\text{SE}(3)$:

$$\xi_1 \circ \xi_2 := \log(\exp(\widehat{\xi}_1) \cdot \exp(\widehat{\xi}_2))^\vee \neq \xi_1 + \xi_2$$
$$\neq \xi_2 \circ \xi_1$$

Intuition: Iteratively solve for $\nabla E(\xi) = 0$ by approximating $\nabla E(\xi)$ *linearly*, (i.e., by approximating $E(\xi)$ *quadratically*)

- using **coarse-to-fine** pyramid approach

Direct Image Alignment

$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} \underbrace{(I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2}_{=: r_i^2(\xi)}$$

1. „Linearize“ \mathbf{r} on Manifold around current pose $\xi^{(n)}$:

$$\mathbf{r}(\xi) \approx \underbrace{\mathbf{r}(\xi^{(k)})}_{\mathbf{r}_0 \in R^n} + \underbrace{\frac{\partial \mathbf{r}(\epsilon \circ \xi^{(k)})}{\partial \epsilon}}_{J_{\mathbf{r}} \in R^{n \times 6}} \Big|_{\epsilon=0} \cdot \underbrace{(\xi \circ (\xi^{(k)})^{-1})}_{\delta_{\xi}}$$

2. Solve for $\nabla E(\xi) = 0$

$$E(\xi) = \|\mathbf{r}_0 + J_{\mathbf{r}} \cdot \delta_{\xi}\|_2^2 = \mathbf{r}_0^T \mathbf{r}_0 + 2\delta_{\xi}^T J_{\mathbf{r}}^T \mathbf{r}_0 + \delta_{\xi}^T J_{\mathbf{r}}^T J_{\mathbf{r}} \delta_{\xi}$$

$$\nabla E(\xi) = 2J_{\mathbf{r}}^T \mathbf{r}_0 + 2J_{\mathbf{r}}^T J_{\mathbf{r}} \delta_{\xi} = 0 \quad \Rightarrow \quad \delta_{\xi} = -(J_{\mathbf{r}}^T J_{\mathbf{r}})^{-1} J_{\mathbf{r}}^T \mathbf{r}_0$$

3. Apply $\xi^{(k+1)} = \delta_{\xi} \circ \xi^{(k)}$

4. Iterate (until convergence)

Direct Image Alignment

$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} \underbrace{(I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2}_{=: r_i^2(\xi)}$$

Requires gradient of residual:

$$\frac{\partial r_i(\epsilon \circ \xi^{(k)})}{\partial \epsilon} \Big|_{\epsilon=0} = \frac{1}{z'} \begin{pmatrix} \nabla I_x f_x & \nabla I_y f_y \end{pmatrix} \begin{pmatrix} 1 & 0 & -\frac{x'}{z'} & -\frac{x'y'}{z'} & (z' + \frac{x'^2}{z'}) & -y' \\ 0 & 1 & -\frac{y'}{z'} & -(z' + \frac{y'^2}{z'}) & \frac{x'y'}{z'} & x' \end{pmatrix}$$

= 1x6 row of J_r

with

- $\begin{pmatrix} x' \\ y' \\ z' \end{pmatrix} := R_{\xi^{(k)}} K^{-1} \begin{pmatrix} dp_{i,x} \\ dp_{i,y} \\ d \end{pmatrix} + \mathbf{t}_{\xi^{(k)}} = \text{warped point (before projection)}$
- $f_x, f_y, K = \text{intrinsic camera calibration}$
- $\nabla I_x, \nabla I_y = \text{image gradients}$

$$E(\xi) = \sum_{\mathbf{p}_i \in \Omega_{\text{ref}}} \underbrace{(I_{\text{ref}}(\mathbf{p}_i) - I(\omega(\mathbf{p}_i, D_{\text{ref}}(\mathbf{p}_i), \xi)))^2}_{=: r_i^2(\xi)}$$

Coarse-to-Fine:

- Minimize for down-scaled image (e.g. factor 8, 4, 2, 1) and use result as initialization for next finer level.
- Elegant formulation:

Downscale image and adjust K correspondingly:

- Downscale by factor of 2 (e.g. 640x480 -> 320->240)

$$\bullet \quad K = \begin{pmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{pmatrix} \quad \rightarrow \quad K_{\frac{1}{2}} = \begin{pmatrix} \frac{f_x}{2} & 0 & \frac{c_x+0.5}{2} - 0.5 \\ 0 & \frac{f_y}{2} & \frac{c_y+0.5}{2} - 0.5 \\ 0 & 0 & 1 \end{pmatrix}$$

- (assuming discrete pixel (x,y) contains continuous value at (x,y))



Final Algorithm:

$$\xi^{(0)} = 0$$

$$k = 0$$

for *level* = 3 ... 1

 compute down-scaled images & depthmaps (factor = 2^{level})

 compute down-scaled K (factor = 2^{level})

for *i* = 1..20

 compute Jacobian $J_r \in R^{n \times 6}$

 compute update δ_ξ

 apply update $\xi^{(k+1)} = \delta_\xi \circ \xi^{(k)}$

k++; maybe break early if δ_ξ too small or if residual increased

done

done

+ robust weights (e.g. Huber), see *iteratively reweighted least squares*

+ *Levenberg-Marquard (LM)* Algorithm

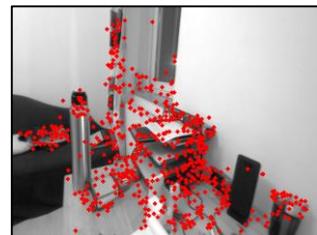
So, what's new?

Keypoint-Based

**Input
Images**



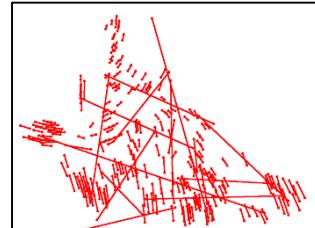
**Extract & Match
Features**
(SIFT / SURF / ...)



abstract images to feature observations

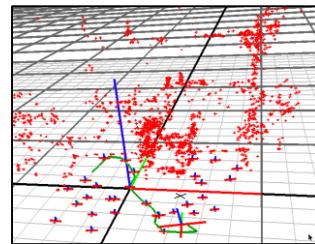
Track:

min. **reprojection** error
(point distances)



Map:

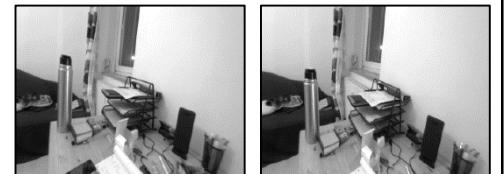
est. **feature-parameters**
(3D points / normals)



can only use & reconstruct corners

Semi-Dense (direct)

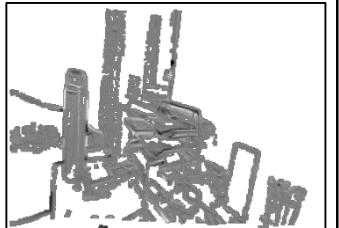
**Input
Images**



keep full image

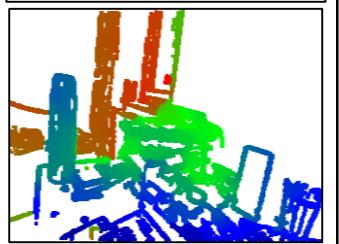
Track:

min. **photometric** error
(intensity difference)



Map:

est. **per-pixel depth**
(semi-dense depth map)



can use & reconstruct whole image