

Robust Odometry Estimation for RGB-D Cameras

Oskar Carlbaum
Guillermo Gonzalez de Garibay

Technische Universität München
Department of Informatics
Computer Vision Group

April 13, 2016

Outline

- 1 Introduction**
- 2 Related Work**
- 3 Main Section**
- 4 Results**
- 5 Conclusions and Improvements**
- 6 Questions/Demo**

Outline

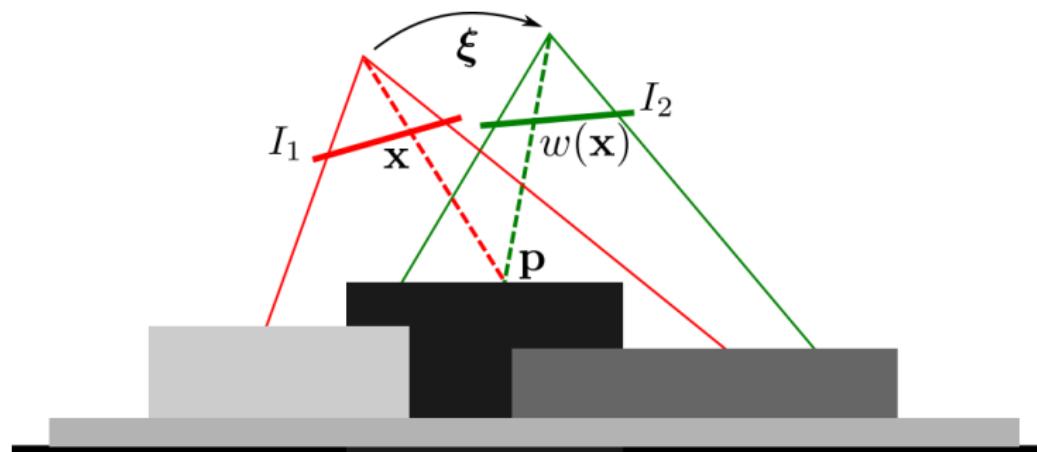
- 1 Introduction
- 2 Related Work
- 3 Main Section
- 4 Results
- 5 Conclusions and Improvements
- 6 Questions/Demo

Introduction

Visual Odometry Estimation

- Estimating camera movements from image sequence
- Uses: Robots, Quadcopters, Mars rover
 - Challenges: Requires high frame rates, robustness to outliers
- Sparse vs Dense
 - Sparse: Feature based detectors for tracking (common)
 - Dense: All image information (more robust)
 - RGB-D images

RGB-D Frames

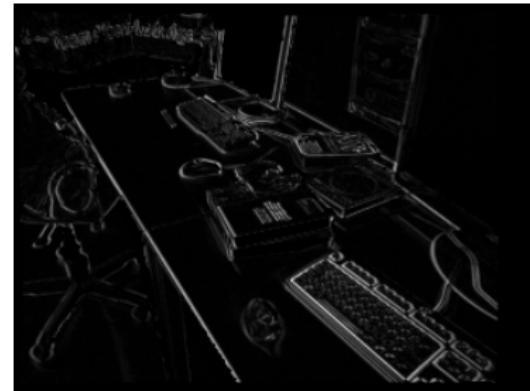


[Kerl et al., 2013]

- Calculate ξ_i^{i+1} for consecutive frames
- Depth information from first frame only
- Align gray images
- Minimize photometric error → See next slide

Photometric error

- $E(\xi) = \sum_i w(r_i)(r_i(\xi))^2$ is to be minimized for consecutive frames
- $r_i(\xi) := I_2(\tau(\xi, \mathbf{x}_i)) - I_1(\mathbf{x}_i)$ (intensity difference after transformation)
 - I_2 not evaluated at grid points \rightarrow __texture__ interpolation
- By linearizing the residuals:
$$\mathbf{r}_{lin}(\xi) = \mathbf{r}(0) + J\Delta\xi$$
- We get to this linear system:
$$J^T W J \Delta\xi = -J^T W \mathbf{r}(0)$$



Twist coordinates

- ξ is just a **minimal** and **smooth** representation of a Rigid Body Motion
- Rigid Body Motion \equiv Rotation + Translation (6DoF)
- Many possible parametrizations (≥ 6 parameters)
 - Most have mathematical singularities
 - \Rightarrow Use twist coordinates: $\xi = (v_1, v_2, v_3, w_1, w_2, w_3)$
- $\xi \longleftrightarrow (R, T)$ have closed-form expressions

Reconstruction

[Kerl et al., 2013]

- Accumulate ξ_i^{i+1} to get ξ_0^t
- Translation is represented above

Outline

- 1 Introduction
- 2 Related Work
- 3 Main Section
- 4 Results
- 5 Conclusions and Improvements
- 6 Questions/Demo

Related work

- Paper:
 - Robust Odometry Estimation for RGB-D Cameras [Kerl et al., 2013]
- Previous implementations:
 - MATLAB: Multiple View Geometry lecture, Exercise Sheet 8 [2015]
 - C++/CUDA: This same course in SS14, WS14/15, SS15

Outline

- 1 Introduction
- 2 Related Work
- 3 Main Section
- 4 Results
- 5 Conclusions and Improvements
- 6 Questions/Demo

Our mission

Goals

- Ludicrous speed
- Code readability
- Fewest libraries
- Robustness
- Comparable results

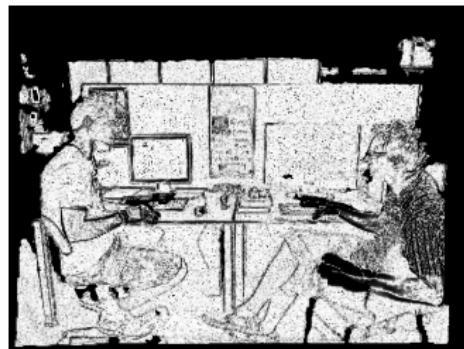
Test material

- TUM Benchmark Datasets
 - RGB-images
 - Depth-images
 - Intrinsic Camera Parameters

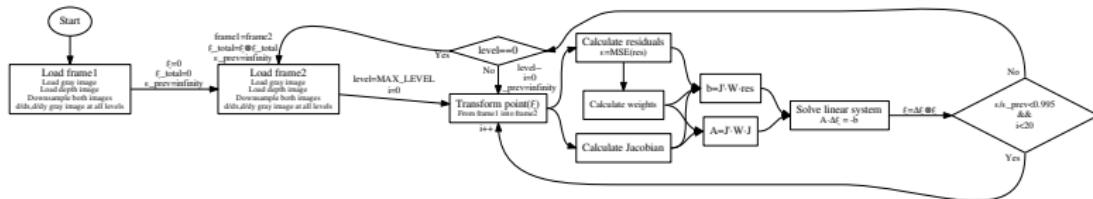
Robustness

Robustness → Weighting

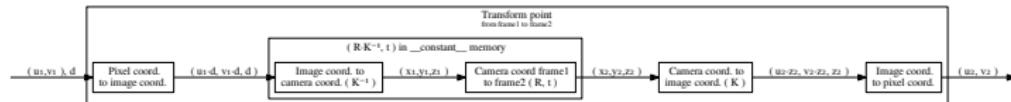
- Movement, blur, reflections, intensity changes
- Student-t distribution → Exclude unreliable data



Flow chart



- Gauss-Newton algorithm
- Gray image and derivatives into texture memory
 - Interpolation for residuals and Jacobian



- The transformation $(R \cdot K^{-1}, t)$ is loaded into constant memory
 - $\sim 10^6$ reads

CUDA

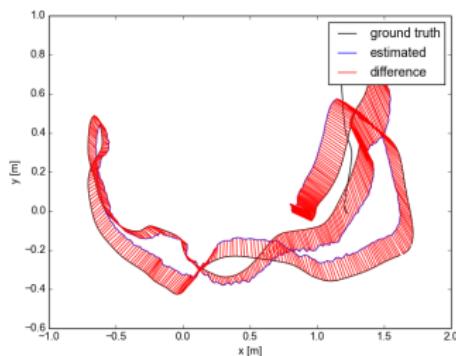
Operations using CUDA:

- Image downscaling
- Intensity image derivatives
- Transform points
- Residuals
- Error → cuBLAS (reduction)
- Weights → cuBLAS (reduction)
- Jacobian
- $J^T W J, J^T W \mathbf{r} \rightarrow$ cuBLAS (matrix mult.)

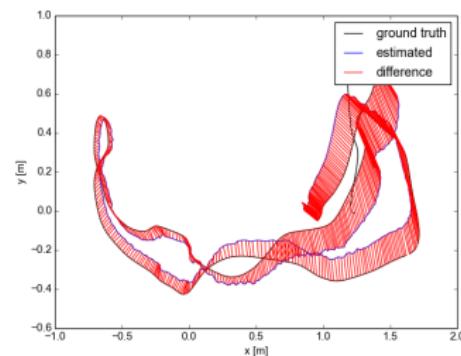
Outline

- 1 Introduction
- 2 Related Work
- 3 Main Section
- 4 Results
- 5 Conclusions and Improvements
- 6 Questions/Demo

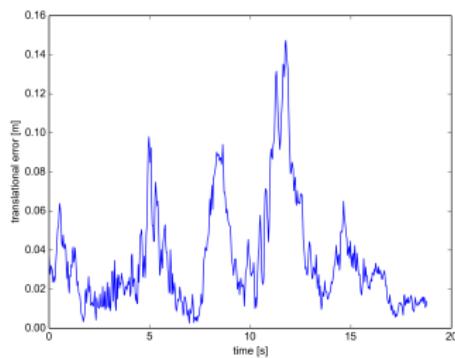
Comparisons - ATE, Desk1



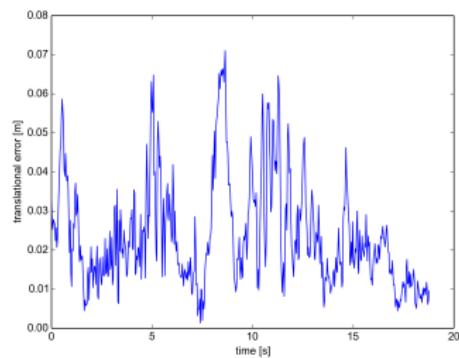
No weights

T-dist. weights
Absolute Trajectory Error

Comparisons - RPE, Desk1



No weights



T-dist. weights
Relative Pose Error

Results

■ Accuracy, TUM-benchmark toolset

Method	fr1/desk	fr1/desk2
Paper: no weights	0.0551	0.1003
Paper: t-dist. weights	0.0458	0.0708
Our: no weights	0.0449	0.0501
Our: t-dist. weights	0.0278	0.0433

Note: we take the translational_error.rmse value from the online_evaluation for relative pose error for pose pairs with a distance of 1 second. We are not completely sure if it is the same value as used in the paper.

■ Speed - runtime

- 'Homemade' method
 - ~ 53.3 ms / frame (no weights)
 - ~ 93.4 ms / frame (t-Dist. weights)
- cuBLAS
 - ~ 26.2 ms / frame (no weights)
 - ~ 44.7 ms / frame (t-Dist. weights)

Outline

- 1 Introduction
- 2 Related Work
- 3 Main Section
- 4 Results
- 5 Conclusions and Improvements
- 6 Questions/Demo

Conclusions

- Works as expected
- Libraries:
 - cuBLAS → easy, fast
 - Eigen → easy CPU matrix handling
 - Sophus (Lie algebra) → replace by 2 functions
- Gaussian for downscaling → not noticeable
- $J^T W J$ just at the highest resolution level, takes about half of the execution time, even with cuBLAS

Improvements

- Motion Prior → Velocity Gaussian distribution
 - Constant velocity assumption
 - Inertial sensors
 - ...
- SLAM
 - Robustness against moving objects
- Do not recalculate Jacobian every iteration (compositional image alignment?)

Outline

- 1 Introduction
- 2 Related Work
- 3 Main Section
- 4 Results
- 5 Conclusions and Improvements
- 6 Questions/Demo

THANKS!

Questions?

Bibliography I

[Kerl et al., 2013] Kerl, C., Sturm, J., and Cremers, D. (2013).
[Robust odometry estimation for rgbd cameras.](#)
In *Int. Conf. on Robotics and Automation*.